

■ Assignment 1 – Dockerized Python + PostgreSQL Stack

This project implements **Assignment 1 (ITCS 6190/8190)** using Docker.

It demonstrates how to containerize a simple **PostgreSQL database** and a **Python app** that connects to the DB, performs queries, and writes results to a shared volume.

■ Project Structure

CloudAssignment/

■ app/

■ Dockerfile # Python app image

■ main.py # Python script that queries DB & writes summary

■ db/

■ Dockerfile # PostgreSQL image with init script

■ init.sql # Schema + seed data (20 trips across 5 cities)

■ out/ # Host folder mounted into container (/out)

■ summary.json # Output file (auto-generated by app)

■ compose.yaml # Docker Compose stack definition

■ makefile # Shortcuts: build, run, clean, etc.

■ README.md # Documentation (this file)

■ Database (PostgreSQL)

- Runs PostgreSQL **16** inside a container.
- Uses `init.sql` to **auto-create schema** and **insert seed data** on first start.
- Data model:

```
CREATE TABLE trips (  
  id SERIAL PRIMARY KEY,  
  city TEXT NOT NULL,  
  minutes INT NOT NULL,  
  fare NUMERIC(6,2) NOT NULL  
);
```

- Seeded with **20 trips** across 5 cities (Charlotte, New York, San Francisco, Chicago, Boston).

■ Application (Python)

- Runs inside a `python:3.11-slim` container.
- Uses `[psycopg3](https://www.psycopg.org/psycopg3/)` to connect to PostgreSQL.
- Reads DB connection info from **environment variables** (`DB_HOST`, `DB_PORT`, etc.).

- Performs 3 queries:
 1. Count total number of trips.
 2. Compute **average fare per city**.
 3. Find **top N longest trips** by minutes (default `N=5`, configurable via `APP_TOP_N`).
- Outputs:
 - - Prints a JSON **summary** to stdout.
 - - Saves the same JSON into `out/summary.json` (via mounted volume).

■■ How to Run

Option A – Using `make`

From the project root:

```
make all
```

This will:

1. Stop any running containers.
2. Clean and recreate the `out/` folder.
3. Build images and start the stack.

Option B – Using Docker Compose directly

```
docker compose up --build
```

■ Example Output

When the app runs, it prints a JSON summary:

```
{
  "total_trips": 20,
  "avg_fare_by_city": [
    {"city": "Boston", "avg_fare": 23.28},
    {"city": "Charlotte", "avg_fare": 18.41},
    {"city": "Chicago", "avg_fare": 24.65},
    {"city": "New York", "avg_fare": 24.55},
    {"city": "San Francisco", "avg_fare": 23.93}
  ],
  "top_by_minutes": [
    {"city": "New York", "minutes": 42, "fare": 45.2},
    {"city": "Boston", "minutes": 38, "fare": 36.25},
    {"city": "Charlotte", "minutes": 35, "fare": 32.4},
    {"city": "San Francisco", "minutes": 33, "fare": 35.8},
```

```
{"city": "Chicago", "minutes": 31, "fare": 30.75}
]
```

This exact file is also saved under:
out/summary.json

■ How to Stop & Clean Up

Stop containers & remove volumes

`docker compose down -v`

Or use Makefile

`make clean`

■ Troubleshooting

- **Error: `failed to read dockerfile`**
→ Ensure `db/Dockerfile` and `app/Dockerfile` exist.
- **App says "Waiting for database..."**
→ Normal on startup; the app retries until the DB is healthy.
- **Port conflict on 5432**
→ Stop local Postgres (`brew services stop postgresql`) or map to a different port in `compose.yml`.
- **No `summary.json` generated**
→ Ensure `out/` folder exists and is mounted (`volumes:` section in compose).
- **Network timeout pulling images**
→ Restart Docker Desktop, switch to a stable network, or manually pull images:

`docker pull postgres:16`

`docker pull python:3.11-slim`

■ Reflection

Through this assignment, I learned:

- How to containerize both **database** and **app layers**.
- How to use **Docker Compose** for multi-service orchestration.
- How to use **volumes** to share output between containers and host.
- Importance of **healthchecks** and retry logic to ensure reliable startup.

For future improvement, I would:

- Add unit tests for the Python logic.
- Parameterize city filters and queries via environment variables.
- Push images to a private Docker Hub repo for easier reuse.

■ Submission Notes

- Repo: *(add your GitHub repo URL here)*
- One-page PDF includes:
 - - Build/run commands.
 - - Example summary JSON (stdout + `summary.json`).
 - - Reflection (above).